

```

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import PolynomialFeatures
from sklearn.tree import DecisionTreeRegressor

```

Data Acquiring

```
df=pd.read_csv("../input/motorcycle-dataset/BIKE DETAILS.csv")
```

```
df
```

	seller_type \	name	selling_price	year
0	Individual	Royal Enfield Classic 350	175000	2019
1	Individual	Honda Dio	45000	2017
2	Individual	Royal Enfield Classic Gunmetal Grey	150000	2018
3	Individual	Yamaha Fazer FI V 2.0 [2016-2018]	65000	2015
4	Individual	Yamaha SZ [2013-2014]	20000	2011
...
1056	Individual	Activa 3g	17000	2010
1057	Individual	Honda CB twister	16000	2012
1058	Individual	Bajaj Discover 125	15000	2013
1059	Individual	Honda CB Shine	12000	2009
1060	Individual	Bajaj Pulsar 150	10000	2008

	owner	km_driven	ex_showroom_price
0	1st owner	350	NaN
1	1st owner	5650	NaN
2	1st owner	12000	148114.0
3	1st owner	23000	89643.0
4	2nd owner	21000	NaN
...
1056	1st owner	500000	52000.0
1057	1st owner	33000	51000.0

```

1058  2nd owner      35000      57000.0
1059  1st owner      53000      58000.0
1060  1st owner      92233      75000.0

```

```
[1061 rows x 7 columns]
```

```
df.isnull().sum()
```

```

name                0
selling_price        0
year                0
seller_type         0
owner               0
km_driven           0
ex_showroom_price    435
dtype: int64

```

```
df.describe()
```

	selling_price	year	km_driven	ex_showroom_price
count	1061.000000	1061.000000	1061.000000	6.260000e+02
mean	59638.151744	2013.867107	34359.833176	8.795871e+04
std	56304.291973	4.301191	51623.152702	7.749659e+04
min	5000.000000	1988.000000	350.000000	3.049000e+04
25%	28000.000000	2011.000000	13500.000000	5.485200e+04
50%	45000.000000	2015.000000	25000.000000	7.275250e+04
75%	70000.000000	2017.000000	43000.000000	8.703150e+04
max	760000.000000	2020.000000	880000.000000	1.278000e+06

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1061 entries, 0 to 1060
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                  1061 non-null  object
1   selling_price         1061 non-null  int64
2   year                  1061 non-null  int64
3   seller_type           1061 non-null  object
4   owner                 1061 non-null  object
5   km_driven             1061 non-null  int64
6   ex_showroom_price     626 non-null   float64
dtypes: float64(1), int64(3), object(3)
memory usage: 58.1+ KB

```

Data Cleaning

**** Handling the Missing values

```
df['ex_showroom_price'].isnull().sum()
```

```
435
```

```
df['ex_showroom_price'].fillna(np.round(df['ex_showroom_price'].median  
( ),2),inplace=True)
```

```
df['ex_showroom_price']
```

```
0      72752.5
```

```
1      72752.5
```

```
2     148114.0
```

```
3      89643.0
```

```
4      72752.5
```

```
...
```

```
1056     52000.0
```

```
1057     51000.0
```

```
1058     57000.0
```

```
1059     58000.0
```

```
1060     75000.0
```

```
Name: ex_showroom_price, Length: 1061, dtype: float64
```

```
np.round(df['ex_showroom_price'],2)
```

```
0      72752.5
```

```
1      72752.5
```

```
2     148114.0
```

```
3      89643.0
```

```
4      72752.5
```

```
...
```

```
1056     52000.0
```

```
1057     51000.0
```

```
1058     57000.0
```

```
1059     58000.0
```

```
1060     75000.0
```

```
Name: ex_showroom_price, Length: 1061, dtype: float64
```

```
df['name'].value_counts().head(15)
```

```
Bajaj Pulsar 150      41
```

```
Royal Enfield Classic 350  27
```

```
Honda Activa [2000-2015]  23
```

```
Honda CB Hornet 160R     22
```

```
Bajaj Pulsar 180       20
```

```
Royal Enfield Thunderbird 350  19
```

```
Hero Honda Splendor Plus  17
```

```
Bajaj Discover 125       16
```

```
Yamaha FZ S V 2.0       16
```

```
Honda Activa 3G         15
```

```
Honda Shine            15
```

```
Honda CB Unicorn 150     14
```

```
TVS Apache RTR 160      14
```

```
Bajaj Pulsar 135 LS     13
```

```
Honda Activa 5G                                     12
Name: name, dtype: int64
```

```
df['owner'].value_counts()
```

```
1st owner      924
2nd owner      123
3rd owner       11
4th owner        3
Name: owner, dtype: int64
```

```
df['seller_type'].value_counts()
```

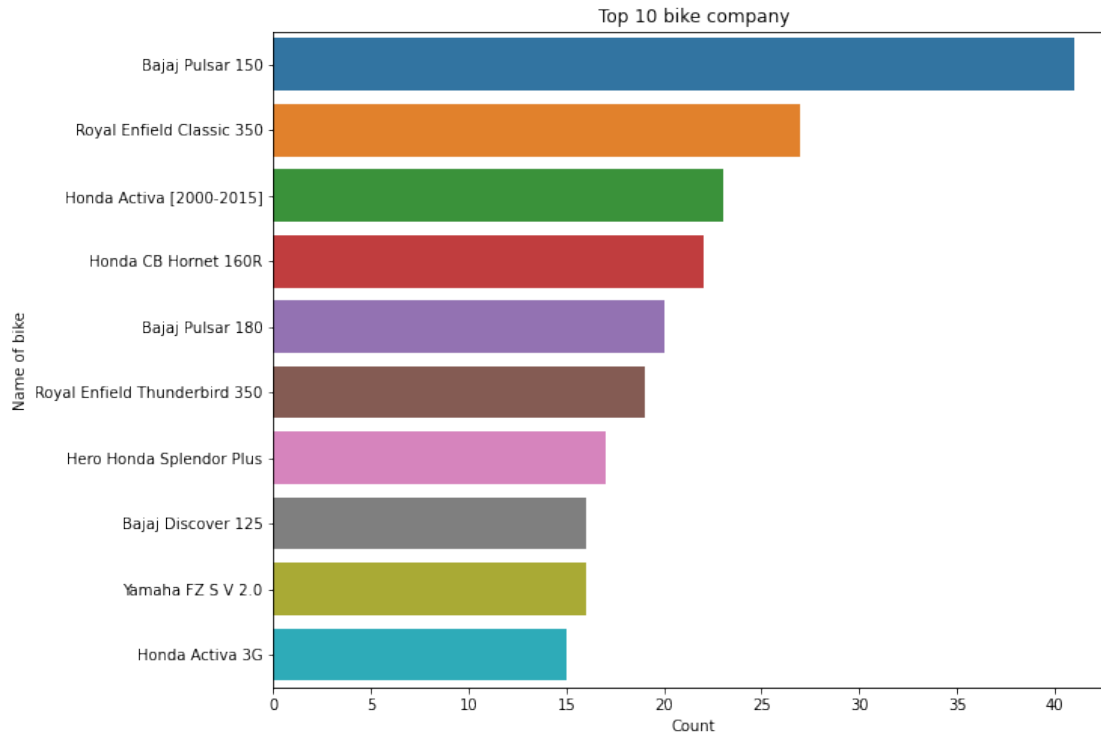
```
Individual      1055
Dealer           6
Name: seller_type, dtype: int64
```

Data Processing

Explorative Data Analysis

```
top_bike_company = df['name'].value_counts().head(10)
plt.figure(figsize=(10, 8))
sns.barplot(x = top_bike_company, y = top_bike_company.index)
plt.ylabel('Name of bike')
plt.title('Top 10 bike company')
plt.xlabel('Count')

Text(0.5, 0, 'Count')
```




```
sns.distplot(df['selling_price'].value_counts())
```


```
/opt/conda/lib/python3.7/site-packages/seaborn/distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed  
in a future version. Please adapt your code to use either `displot` (a  
figure-level function with similar flexibility) or `histplot` (an  
axes-level function for histograms).
```

```
warnings.warn(msg, FutureWarning)
```


```
<AxesSubplot:xlabel='selling_price', ylabel='Density'>
```



```
sns.countplot(x=df['year'].head(10))
<AxesSubplot:xlabel='year', ylabel='count'>
```



```
sns.countplot(x=df['seller_type'], data=df)
<AxesSubplot:xlabel='seller_type', ylabel='count'>
```



Handling the Categorical values

we can convert categorical data to binary data through OneHotEncoder method or get_dummies method

```
df=pd.get_dummies(df,columns=['owner','seller_type'],drop_first=True)
```

'''

or

from sklearn.preprocessing import OneHotEncoder

encoder=OneHotEncoder(drop='first')

enc=pd.DataFrame(encoder.fit_transform(df[['owner']]).toarray())

'''

```
"\nor\nfrom sklearn.preprocessing import OneHotEncoder\
```

```
nencoder=OneHotEncoder(drop='first')\
```

```
nenc=pd.DataFrame(encoder.fit_transform(df[['owner']]).toarray())\n"
```

```
def bike_model(model_name, excl_honda_hero=False):
```

```
    model_list = []
```

```
    if excl_honda_hero:
```

```
        for i in df['name']:
```

```
            if model_name in i and 'Hero' not in i:
```

```
                model_list.append(i)
```

```
    return model_list
```

```
else:
```

```
    for i in df['name']:
```

```
        if model_name in i:
```

```
            model_list.append(i)
```

```
    return model_list
```

```

royal_enfield = bike_model('Royal Enfield')
honda = bike_model('Honda',excl_honda_hero=True)
bajaj = bike_model('Bajaj')
yamaha = bike_model('Yamaha')
suzuki = bike_model('Suzuki')
hero = bike_model('Hero')
tvs = bike_model('TVS')
ktm = bike_model('KTM')

```

```

def brand(i):
    if i in royal_enfield:
        return 'Royal Enfield'
    elif i in honda:
        return 'Honda'
    elif i in bajaj:
        return 'Bajaj'
    elif i in yamaha:
        return 'Yamaha'
    elif i in hero:
        return 'Hero'
    elif i in tvs:
        return 'TVS'
    elif i in suzuki:
        return 'Suzuki'
    elif i in ktm:
        return 'KTM'
    else:
        return 'Other'

```

```

df['brand'] = df['name'].apply(lambda x:brand(x))
df.head()

```

	name	selling_price	year	km_driven
0	Royal Enfield Classic 350	175000	2019	350
1	Honda Dio	45000	2017	5650
2	Royal Enfield Classic Gunmetal Grey	150000	2018	12000
3	Yamaha Fazer FI V 2.0 [2016-2018]	65000	2015	23000
4	Yamaha SZ [2013-2014]	20000	2011	21000

	ex_showroom_price	owner_2nd owner	owner_3rd owner	owner_4th
0	72752.5	0	0	
1	72752.5	0	0	


```

0
2          148114.0          0          0
0
3          89643.0          0          0
0
4          72752.5          1          0
0

```

```

    seller_type_Individual    brand
0                1  Royal Enfield
1                1         Honda
2                1  Royal Enfield
3                1         Yamaha
4                1         Yamaha

```

```
df=df.drop('name',axis='columns')
```

```
df= pd.get_dummies(df, columns=['brand'], drop_first=True)
```

```
x=df.drop('selling_price',axis='columns')
```

```
y=df['selling_price']
```

Train Test Split

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```

Data Modeling

```
lin_reg=LinearRegression()
```

```
lin_reg.fit(x_train,y_train)
```

```
LinearRegression()
```

```
print("Model Intercept: ",lin_reg.intercept_)
```

```
print("Model Coefficient: ",lin_reg.coef_)
```

```
print("Model Coefficient: ",sum(lin_reg.coef_))
```

```
Model Intercept: -7378244.053604637
```

```
Model Coefficient: [ 3.67555304e+03 -4.31636574e-02  5.31133013e-01 -
5.68265614e+03
```

```
 1.97106293e+04  1.78218123e+04 -1.41637480e+04 -3.77752986e+03
```

```
-3.14007480e+03  8.45983695e+04  2.08886295e+04  5.98193684e+04
```

```
-1.08881358e+03  3.21982099e+03  5.13215333e+03]
```

```
Model Coefficient: 187014.0018849024
```

```
y_pred_test=lin_reg.predict(x_test)
```

```
y_pred_train=lin_reg.predict(x_train)
```

```
lin_reg
```

```
LinearRegression()
```

```
lin_reg.predict([[2019,350,72752.5,0,0,0,1,0,0,0,0,1,0,0,0]])
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/base.py:451:
UserWarning: X does not have valid feature names, but LinearRegression
was fitted with feature names
```

```
"X does not have valid feature names, but"
```

```
array([126979.29749026])
```

```
sum(y_pred_train)
```

```
46420892.00000158
```

```
train_res=y_train-y_pred_train
```

```
test_res=y_test-y_pred_test
```

```
train_res
```

```
950    -1056.951098
```

```
782    -1069.378083
```

```
749   -19049.202127
```

```
381   -24245.328807
```

```
958    15516.206691
```

```
...
```

```
337    -4168.860505
```

```
502    17390.550009
```

```
675     1370.210492
```

```
853     9162.222489
```

```
866   -70408.846196
```

```
Name: selling_price, Length: 795, dtype: float64
```

```
test_res
```

```
105     18128.875915
```

```
971   -48972.025308
```

```
461   -3816.790146
```

```
754    22255.880623
```

```
230     1717.231501
```

```
...
```

```
662   -18087.661904
```

```
519     5830.945276
```

```
635     2090.762243
```

```
194   -16176.087586
```

```
916   -15137.524791
```

```
Name: selling_price, Length: 266, dtype: float64
```

```
from sklearn.metrics import
```

```
mean_squared_error,mean_absolute_error,r2_score
```

```
Mean Squared error and Root Mean Squared error
```

```
mse=mean_squared_error(y_test,y_pred_test)
```

```
mse
```

```
1064870473.0227611
```

```

mse_train=mean_squared_error(y_train,y_pred_train)
mse_train

528138316.53659874

rmse=np.sqrt(mse)
rmse

32632.353164041986

rmse=np.sqrt(mse_train)
rmse

22981.2601163774

lin_reg.score(x_train,y_train)

0.818041848370445


lin_reg.score(x_test,y_test)

0.7297034397025874

plt.scatter(y_test,test_res)

<matplotlib.collections.PathCollection at 0x7fc55da0de90>

```



```
sns.distplot(test_res,kde=True)
```


```

/opt/conda/lib/python3.7/site-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a
figure-level function with similar flexibility) or `histplot` (an

```


```
axes-level function for histograms).  
warnings.warn(msg, FutureWarning)
```

```
<AxesSubplot:xlabel='selling_price', ylabel='Density'>
```



```
plt.scatter(y_pred_test, test_res, c='r')  
plt.axhline(y=0, color='blue')
```

```
<matplotlib.lines.Line2D at 0x7fc55da8a9d0>
```



```
import statsmodels.formula.api as smf
modell=smf.ols('y~x',data=df).fit()
modell.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

OLS Regression Results

```
=====
=====
Dep. Variable:                y    R-squared:
0.796
Model:                        OLS    Adj. R-squared:
0.794
Method:                    Least Squares    F-statistic:
272.6
Date:                Thu, 21 Jul 2022    Prob (F-statistic):
0.00
Time:                17:11:10    Log-Likelihood:
-12266.
No. Observations:                1061    AIC:
2.456e+04
Df Residuals:                1045    BIC:
2.464e+04
Df Model:                15

Covariance Type:                nonrobust

=====
=====
```

	coef	std err	t	P> t	[0.025
0.975]					

Intercept	-7.201e+06	4.06e+05	-17.715	0.000	-8e+06
-6.4e+06					
x[0]	3582.7039	201.943	17.741	0.000	3186.444
3978.964					
x[1]	-0.0364	0.017	-2.179	0.030	-0.069
-0.004					
x[2]	0.5959	0.015	40.585	0.000	0.567
0.625					
x[3]	-5577.7978	2512.229	-2.220	0.027	-1.05e+04
-648.211					
x[4]	2.029e+04	7915.992	2.563	0.011	4754.873
3.58e+04					
x[5]	6203.1532	1.57e+04	0.395	0.693	-2.46e+04
3.7e+04					
x[6]	-8504.5739	1.05e+04	-0.808	0.419	-2.91e+04
1.21e+04					
x[7]	-5210.9224	2322.398	-2.244	0.025	-9768.016
-653.828					
x[8]	-3935.8774	2418.645	-1.627	0.104	-8681.831
810.076					
x[9]	7.6e+04	5529.175	13.745	0.000	6.51e+04
8.68e+04					
x[10]	1.672e+04	5289.319	3.160	0.002	6337.309
2.71e+04					
x[11]	5.675e+04	2982.476	19.029	0.000	5.09e+04
6.26e+04					
x[12]	7952.5339	4961.751	1.603	0.109	-1783.595
1.77e+04					
x[13]	3307.3299	3476.214	0.951	0.342	-3513.825
1.01e+04					
x[14]	2101.5975	3023.855	0.695	0.487	-3831.922
8035.117					
=====					
=====					
Omnibus:		728.287	Durbin-Watson:		
1.730					
Prob(Omnibus):		0.000	Jarque-Bera (JB):		
23729.204					
Skew:		2.687	Prob(JB):		
0.00					
Kurtosis:		25.536	Cond. No.		
5.47e+07					
=====					
=====					

Notes:

```
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
[2] The condition number is large, 5.47e+07. This might indicate that
there are
strong multicollinearity or other numerical problems.
"""
```

Saving model and Converting

```
from joblib import dump, load
```

```
dump(lin_reg, "motor_cycle_prediction.joblib")
```

```
['motor_cycle_prediction.joblib']
```

#Deployment

```
loaded_lin=load("motor_cycle_prediction.joblib")
```

```
loaded_model=load("motor_cycle_prediction.joblib")
```

```
loaded_model.predict([[2019,350,72752.5,0,0,0,1,0,0,0,0,1,0,0,0]])
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/base.py:451:
```

```
UserWarning: X does not have valid feature names, but LinearRegression
was fitted with feature names
```

```
    "X does not have valid feature names, but"
```

```
array([126979.29749026])
```